

CS 486 — Lecture 5: Local Search

1 Intro to Local Search

- We've discussed a few search strategies — but they have some problems.
- So far, search algorithms explore the space systematically and keep track of one or more paths.
- What happens if the search space is too big?
- What if we don't care about the path to the goal — just the goal node itself? For example, nobody cares about the *order* of the solution to the 4Q problem. Just the result.
- We remove two things we care about. One is that we no longer systematically search. Furthermore, we no longer remember the path to the goal node — just the current state we are visiting.
- This means we need less memory and only explore a portion of the search space. It can find solutions quickly on average and works on CSPs and general and general optimization problems.
- However, there is no guarantee that a solution will be found if one exists, nor can we prove that no solution exists with just a local search.
- A LSP contains of:
 - A state — a complete assignment to all variables.
 - A neighbour relation — which states to explore next?
 - A cost function — how good is each state?
- Let us show the 4Q problem as a LSP:
 - Variables and domains are the same as before.
 - The initial state is 4 queens on the board in random row positions.
 - The goal state is 4 queens on the board that are not attacking each other.
 - Note, we have no constraints — this is as there are many possible positions that we check that will not meet the constraints.
 - The neighbour relation can either be moving a single queen to a different row or swap the row positions of 2 queens. Discussed later.
 - The cost function is that we want to minimize the number of pairs of queens attacking each other, directly or indirectly (so if you had 3 queens on the diagonal, then the first and third queen are technically indirectly attacking each other).

2 Local Search Algorithms

2.1 Greedy Descent

- Start with a random state.
- Move to a neighbour with the lowest cost if it is better than the current state.
- Stop when no neighbour has a lower cost than the current state. It will be one of the best states in the neighbourhood.
- Greedy as it only sees and considers the current node (a la greedy algorithm).

- This performs well and will often progress rapidly towards a solution.
- However, greedy descent can get stuck at local optimums (think of the gradient descent pits).
- GD can get stuck in actual local minimums or flat local minimums — for example, a long flat section that isn't a pit (shoulder), or an actual flat pit. The latter is much harder to escape, the former can get escaped if we travel long enough.
- Ways we can improve is allowing moving to states with the same cost, and keeping some track of recent neighbours to avoid getting stuck in loops.
- One way to deal with getting stuck is random restarts (restart randomly in a new location) and random walks (don't always move to a state with a lower cost).

2.2 Simulated Annealing

- GD focuses on optimization/exploitation, whereas randomly moving allows us to explore the search space.
- Can we combine these two properties into one algorithm?
- Annealing is a process where we cool down molten metals to make them stronger.
- As per CS tradition, we stole this term to describe what we want to do.
- At each step, choose a random neighbour. If it is an improvement, move to it. If it is not, then move to the neighbour probabilistically depending on both the “current temperature”, T , and how much worse the neighbour is compared to the compared state.
- Define A as the current state and A' as the worst neighbour. Then let $\Delta C = \text{cost}(A') - \text{cost}(A)$. The current temperature is T .
- Then, our probability of moving to neighbour A' is $e^{-\frac{\Delta C}{T}}$.
- We can write the simulated annealing function as such:

```
current = initial_state
T = some large positive value
while T > 0:
    next = a random neighbour of current
    delta_c = cost(next) - cost(current)
    if delta_c < 0:
        current = next
    else:
        current = next with prob. p = e**(-delta_c / T)
    decrease T
return current
```

- In theory, we want to decrease T slowly to find the global optimum. If it is slow enough, we are guaranteed to find the global optimum with prob. approaching 1.
- An “annealing schedule” that's commonly used is called geometric cooling (ie: multiply T by 0.99 each step).

2.3 Genetic Algorithms

- Local search algorithms so far only remember one state — what if we remember multiple states?
- One idea is beam search.
 - Remember k states, and choose the k best states out of all neighbours.

- If $k = 1$, then this is identical to a normal local search.
- How is beam search different from k random restarts in parallel? The idea behind beam searching is that it is aware of other states — restarting randomly would not have this ability, and are independent!
- But this communication adds a problem — now you have a good region in the search space, then we cluster our results around there, losing some diversity.
- We next look at stochastic beam search. Beam search chooses k states deterministically; stochastic BS chooses via probabilities.
 - The probability of choosing a neighbour is prop. to its fitness.
 - This maintains diversity in the population of states.
 - This also mimics natural selection.
- Lastly, we have the genetic algorithm. Similar to the stochastic BS, but it differs in that we pick two states to make a new child state, and there is a random probability for the child state to mutate.