

CS 486 — Lecture 6: Machine Learning

1 Introduction

- Some applications:
 - Medical diagnosis
 - Spam filtering
 - Facial recognition
 - Speech understanding
 - Handwriting recognition
- Learning is the ability of an agent to improve its performance on future tasks based on experience.
- So the things we want an agent to do are to do more (expand range of behaviours), do things better (improve accuracy), and do things faster.
- Why learn over just hardcoding?
 - There may be situations which cannot be anticipated.
 - Situations may change over time which cannot be anticipated.
 - Sometimes, some solutions just can't be programmed — how do you program image recognition, for example?
- Any learning problem consists of these parts:
 - The problem/task
 - Experience/data that we are using to improve the performance of our agent
 - Background knowledge/bias (make assumptions about the world/task)
 - Measure of improvement — how do we know if we are doing better/worse?
- Types of learning:
 - Supervised learning — given input features, target features, and training examples, predict the value of the target features for new examples given their values on the input features.
 - Unsupervised learning — learning classifications when the examples do not have the targets defined.
 - Reinforcement learning — learning what to do based on rewards and punishments.
- Let us focus on two types of supervised learning problems:
 1. Classification — target features are discrete (ie: is the weather tomorrow sunny, cloudy, or rainy).
 2. Regression — target features are continuous (ie: tomorrow's temperature).

2 Supervised Learning

- Given training examples of the form $(x, f(x))$, we want to return a function h (hypothesis) that approximates f .
- This does assume that some true function $f(x)$ exists; but during training we never get to actually see f (otherwise we would just use it).

- We can see learning as a search problem if we see it as looking for one best hypothesis given a hypothesis “space”.
- The search space is often too large for a systematic search, so most ML techniques are some form of a local search.
- The goal of ML is to find a hypothesis that can predict unseen examples correctly. If it can do so well, then we say it generalizes well.
- How can we choose a hypothesis that generalizes well?
- Occam’s razor — assume the simplest hypothesis is the best.
- Cross-validation — use a test set later to validate your hypothesis after using the training set on it.
- There is often a trade-off between complex hypotheses that fit the training data better, and simpler hypotheses that generalize better but fit worse.
- The bias-variance trade-off — error is at its highest when the model is too complex (bias low, variance high) OR if it is too low (bias high, variance low) (so a U shape). We want to find a sweet spot in the middle.
- The bias is about the effect of the hypothesis on how well we can fit the data.
- A hypothesis with high bias will be too simplistic, have too few degrees of freedom, assumes too much, and does not fit the training data well.
- If there is too high bias, it cannot use the training data well.
- Meanwhile, the variance is about how much the learned hypothesis varies given different training data.
- A hypothesis with high variance has many degrees of freedom, is very flexible, fits the training data very well, and is very sensitive — small changes in the training data can cause the hypothesis to change a lot.
- This also means that it cannot adapt to test data; it is overfitted to training data and cannot generalize.